

Exact Scheduling of Parallel Additive Manufacturing Machines for Total Tardiness Minimization

Abstract—This paper studies the one-dimensional batch scheduling problem for parallel additive manufacturing machines with the objective of minimizing total tardiness. To solve this problem exactly, we develop tailored branch-and-bound algorithm.

I. MATHEMATICAL PROGRAMMING MODEL

A. Problem Description

A set of parts $J = \{1, \dots, n\}$ is to be processed on a set of identical parallel metal additive manufacturing machines $M = \{1, \dots, \bar{M}\}$ based on Powder Bed Fusion technology. Without loss of generality, we assume that each customer order consists of a single part, so that order tardiness is equivalent to part tardiness. Each part $j \in J$ has a due date d_j , a volume v_j , and fixed geometric dimensions given by length l_j , width w_j , and height h_j . All machines are assumed to have identical rectangular build platforms, each with length L and width W .

Parts are assigned to machines and grouped into batches. On each machine, batches are processed sequentially from time zero, while different machines can process their assigned batches in parallel. We consider a one-dimensional capacity restriction, under which a set of parts can be assigned to the same batch if the sum of their projected areas does not exceed the platform capacity $L \times W$. This assumption gives rise to the identical parallel-machine one-dimensional batch scheduling problem studied in this paper.

The processing time of a batch consists of three components: the fixed machine setup time S , the laser scanning time, and the recoating time [7]. More specifically, the laser scanning time depends on the total volume of the parts assigned to the batch, whereas the recoating time depends on the maximum height of the parts in the batch. Let P_{mb} denote the processing time of batch position b on machine m , and let C_{mb} denote its completion time. If part j is assigned to batch position b on machine m , then its completion time c_j is equal to C_{mb} . The tardiness of part j is defined as

$$T_j = \max\{0, c_j - d_j\}.$$

All parts are assumed to be available at time zero. The objective is to determine the assignment of parts to machines, the grouping of parts into batches on each machine, and the

sequence of batches on each machine so as to minimize the total tardiness of all parts.

B. MILP Formulation

We formulate the identical parallel-machine one-dimensional batch scheduling problem as a mixed-integer linear program (MILP). The formulation extends the single-machine model by introducing machine-indexed batch positions. Each part must be assigned to exactly one batch on exactly one machine, and batches on the same machine are processed sequentially from time zero. Since the machines are identical, the processing-time structure of each batch remains the same across machines.

Sets and Indices

- $J = \{1, \dots, n\}$: Set of parts.
- $M = \{1, \dots, \bar{M}\}$: Set of identical parallel machines.
- $B = \{1, \dots, n\}$: Set of available batch positions on each machine.

Parameters

- l_j, w_j, h_j : Length, width, and height of part j .
- v_j : Volume of part j .
- d_j : Due date of part j .
- L, W : Length and width of the build platform.
- S : Setup time of a machine.
- V : Laser scanning time coefficient.
- U : Recoating time coefficient.
- $\mathcal{M}$: A sufficiently large positive constant.

Decision Variables

- x_{jmb} : equal to 1 if part j is assigned to batch position b on machine m , 0 otherwise.
- z_{mb} : equal to 1 if batch position b on machine m is used, 0 otherwise.
- H_{mb} : height of batch b on machine m .
- P_{mb} : processing time of batch b on machine m .
- C_{mb} : completion time of batch b on machine m .
- c_j : completion time of part j .
- T_j : tardiness of part j .

The MILP formulation is given as follows:

$$\text{Minimize} \quad \sum_{j \in J} T_j \quad (1)$$

s.t.

This paragraph of the first footnote will contain the author affiliations, contact information, and email addresses in the final camera-ready version. It is currently blinded to comply with double-anonymous review guidelines.

This paragraph of the second footnote will contain funding information and grant numbers in the final camera-ready version. It is currently blinded to comply with double-anonymous review guidelines.

$$\sum_{m \in M} \sum_{b \in B} x_{jmb} = 1, \quad \forall j \in J \quad (2)$$

$$\sum_{j \in J} x_{jmb} \leq M z_{mb}, \quad \forall m \in M, \forall b \in B \quad (3)$$

$$\sum_{j \in J} (l_j w_j) x_{jmb} \leq LW, \quad \forall m \in M, \forall b \in B \quad (4)$$

$$H_{mb} \geq h_j x_{jmb}, \quad \forall j \in J, \forall m \in M, \forall b \in B \quad (5)$$

$$P_{mb} = S z_{mb} + V \sum_{j \in J} v_j x_{jmb} + U H_{mb}, \quad \forall m \in M, \forall b \in B \quad (6)$$

$$C_{m1} \geq P_{m1}, \quad \forall m \in M \quad (7)$$

$$C_{mb} \geq C_{m,b-1} + P_{mb}, \quad \forall m \in M, \forall b = 2, \dots, |B| \quad (8)$$

$$c_j \geq C_{mb} - M(1 - x_{jmb}), \quad \forall j \in J, \forall m \in M, \forall b \in B \quad (9)$$

$$T_j \geq c_j - d_j, \quad \forall j \in J \quad (10)$$

$$T_j \geq 0, \quad \forall j \in J \quad (11)$$

$$z_{m,b+1} \leq z_{mb}, \quad \forall m \in M, \forall b = 1, \dots, |B| - 1 \quad (12)$$

$$x_{jmb}, z_{mb} \in \{0, 1\}, \quad H_{mb}, P_{mb}, C_{mb}, c_j, T_j \geq 0 \quad (13)$$

The objective function (1) minimizes the total tardiness of all parts. Constraints (2) ensure that each part is assigned to exactly one batch on exactly one machine. Constraints (3) activate batch position b on machine m whenever at least one part is assigned to it. Constraints (4) enforce the one-dimensional build-platform capacity restriction for every batch. Constraints (5) define the batch height as the maximum height of the parts assigned to that batch. Constraints (6) calculate the processing time of each batch. Constraints (7)–(8) define the sequential completion times of batches on each machine. Constraint (9) links each part completion time to the completion time of its assigned batch. Constraints (10)–(11) define part tardiness. Finally, constraints (12) enforce contiguous batch usage on each machine, i.e., batch position $b + 1$ on machine m cannot be used unless batch position b on the same machine is used.

Because the machines are identical, the model also contains symmetry across machines. To mitigate this effect, an additional symmetry-breaking constraint may be introduced:

$$\sum_{b \in B} z_{mb} \geq \sum_{b \in B} z_{m+1,b}, \quad \forall m = 1, \dots, \bar{M} - 1 \quad (14)$$

Constraint (14) imposes a nonincreasing order on the number of used batch positions across machines. In other words, machine $(m+1)$ cannot use more batch positions than machine (m) . Since the machines are identical, this restriction does not exclude any structurally distinct feasible schedule, but it helps reduce redundant symmetric solutions and can improve the tractability of the MILP.

II. METHODOLOGY

A. Overall Branch-and-Bound Framework

This section presents the overall exact solution framework for the identical parallel-machine AM batch scheduling

problem. The proposed method follows a decomposition-oriented branch-and-bound structure. The upper-level search tree branches on part-to-machine assignments, while the detailed batching and sequencing decisions on each machine are evaluated by an exact single-machine scheduling oracle.

At the beginning of the algorithm, an initial feasible solution is constructed by a fast heuristic, which provides the initial global upper bound UB . The search is initialized at the root node N^0 , where no part has been assigned to any machine:

$$\mathcal{S}_m(N^0) = \emptyset, \quad \forall m \in \mathcal{M},$$

and

$$\mathcal{U}(N^0) = \mathcal{J}.$$

The root node is inserted into the active node list.

During the search, an active node N is selected according to a node selection rule. The algorithm first evaluates a valid lower bound $LB(N)$. If

$$LB(N) \geq UB,$$

then node N cannot lead to a solution better than the incumbent and is pruned. Otherwise, the algorithm checks whether N represents a complete assignment. If

$$\mathcal{U}(N) = \emptyset,$$

then every part has been assigned to one machine. In this case, the problem decomposes into $|\mathcal{M}|$ independent single-machine scheduling subproblems. For each machine m , the exact oracle is called to solve the subproblem induced by the assigned part set $\mathcal{S}_m(N)$. The objective value of the complete assignment is then computed as

$$Z(N) = \sum_{m \in \mathcal{M}} \Phi(\mathcal{S}_m(N)),$$

where $\Phi(\mathcal{S}_m(N))$ denotes the optimal total tardiness of the single-machine problem associated with $\mathcal{S}_m(N)$. If $Z(N) < UB$, the incumbent solution and the global upper bound are updated.

If N is not a complete assignment and satisfies $LB(N) < UB$, the node is branched by selecting one unassigned part $j \in \mathcal{U}(N)$ and assigning it to one of the parallel machines. Each resulting child node fixes one additional part-to-machine assignment. For each generated child node, the lower bound is evaluated, and the child node is inserted into the active node list only if its lower bound is smaller than the current upper bound.

The algorithm terminates when the active node list becomes empty. At this point, every unexplored node has either been evaluated exactly or pruned by a valid lower bound. Therefore, the incumbent solution is guaranteed to be globally optimal.

The overall procedure is summarized as follows.

- 1) Construct an initial feasible solution and set its objective value as the initial upper bound UB .
- 2) Initialize the root node N^0 with $\mathcal{S}_m(N^0) = \emptyset$ for all $m \in \mathcal{M}$ and $\mathcal{U}(N^0) = \mathcal{J}$. Insert N^0 into the active node list $\mathcal{L}$.

- 3) While $\mathcal{L}$ is not empty, select a node N from $\mathcal{L}$ according to the node selection rule.
- 4) Compute a valid lower bound $LB(N)$. If $LB(N) \geq UB$, prune node N .
- 5) If $\mathcal{U}(N) = \emptyset$, evaluate the complete assignment by solving the single-machine subproblem induced by each $\mathcal{S}_m(N)$ using the exact oracle. Compute

$$Z(N) = \sum_{m \in \mathcal{M}} \Phi(\mathcal{S}_m(N)).$$

If $Z(N) < UB$, update the incumbent solution and set $UB \leftarrow Z(N)$.

- 6) If $\mathcal{U}(N) \neq \emptyset$ and $LB(N) < UB$, select a branching part $j \in \mathcal{U}(N)$ and generate child nodes by assigning j to candidate machines.
- 7) For each generated child node N' , compute or update $LB(N')$. If $LB(N') < UB$, insert N' into $\mathcal{L}$; otherwise, prune N' .
- 8) When $\mathcal{L}$ becomes empty, return the incumbent solution and its objective value UB .

B. Assignment-Based Search Tree

The upper-level search tree is constructed over part-to-machine assignments. Instead of explicitly branching on batch compositions and batch sequences, these machine-internal decisions are deferred to the single-machine oracle. Therefore, each node in the upper-level tree records only which parts have already been fixed to each machine and which parts remain unassigned.

1) *Node Representation:* Let $\mathcal{M} = \{1, \dots, M\}$ denote the set of identical parallel machines and let $\mathcal{J} = \{1, \dots, n\}$ denote the set of parts. For a node N , let

$$\mathcal{S}_m(N) \subseteq \mathcal{J}, \quad m \in \mathcal{M},$$

be the set of parts that have already been assigned to machine m . Let

$$\mathcal{U}(N) = \mathcal{J} \setminus \bigcup_{m \in \mathcal{M}} \mathcal{S}_m(N)$$

be the set of parts that remain unassigned at node N .

Thus, node N is represented as

$$N = (\mathcal{S}_1(N), \dots, \mathcal{S}_M(N), \mathcal{U}(N)),$$

where the assigned sets are mutually disjoint, i.e.,

$$\mathcal{S}_m(N) \cap \mathcal{S}_{m'}(N) = \emptyset, \quad \forall m \neq m'.$$

The root node N^0 is defined by

$$\mathcal{S}_m(N^0) = \emptyset, \quad \forall m \in \mathcal{M}, \quad \mathcal{U}(N^0) = \mathcal{J}.$$

A node is called terminal if all parts have been assigned, namely $\mathcal{U}(N) = \emptyset$. At a terminal node, the parallel-machine problem decomposes into M independent single-machine subproblems. The exact single-machine oracle is then called for each assigned set $\mathcal{S}_m(N)$ to evaluate the complete assignment.

2) *Symmetry-Reduced Branching Scheme:* At a nonterminal node N , one unassigned part $j \in \mathcal{U}(N)$ is selected for branching. A straightforward branching rule would assign part j to any machine $m \in \mathcal{M}$. However, because the machines are identical, this rule creates many symmetric nodes. For example, assigning the first part to machine 1 or machine 2 yields equivalent partial assignments up to a relabeling of machines.

To reduce such symmetry, we impose a canonical machine-opening rule. At any node N , let

$$q(N) = |\{m \in \mathcal{M} : \mathcal{S}_m(N) \neq \emptyset\}|$$

denote the number of machines that have already been used. Machines are opened in their index order. Therefore, the selected part j can be assigned either to one of the already used machines or to the first unused machine. The candidate machine set is defined as

$$\mathcal{K}(N) = \begin{cases} \{1\}, & q(N) = 0, \\ \{1, \dots, \min\{q(N) + 1, M\}\}, & q(N) \geq 1. \end{cases}$$

For each candidate machine $r \in \mathcal{K}(N)$, a child node $N^{(r)}$ is generated by assigning part j to machine r :

$$\mathcal{S}_r(N^{(r)}) = \mathcal{S}_r(N) \cup \{j\},$$

$$\mathcal{S}_m(N^{(r)}) = \mathcal{S}_m(N), \quad \forall m \in \mathcal{M}, m \neq r,$$

and

$$\mathcal{U}(N^{(r)}) = \mathcal{U}(N) \setminus \{j\}.$$

This branching rule preserves completeness. Since the machines are identical, any feasible assignment that uses a later unused machine before an earlier unused machine has an equivalent assignment obtained by renaming machines. Therefore, restricting the search to the canonical machine-opening order does not remove any structurally distinct feasible solution. It only eliminates redundant symmetric representations.

3) *Branching Part Selection:* The efficiency of the assignment-based search depends strongly on the choice of the part selected for branching. A simple rule is to select the unassigned part with the earliest due date, since such parts are usually more critical for the total tardiness objective. Another possible rule is to select the part with the largest projected area, volume, or height, because such parts have a stronger impact on batching feasibility and processing times.

In this study, we use a lower-bound-based branching score to select the branching part. The idea is to choose a part whose assignment to different machines leads to the largest variation in the lower bounds of the child nodes. Such a part is more informative for the search because it can better differentiate promising and unpromising assignment decisions.

For each candidate part $j \in \mathcal{U}(N)$ and each candidate machine $r \in \mathcal{K}(N)$, let $N_j^{(r)}$ denote the child node obtained by assigning j to machine r . A tentative lower bound

$LB(N_j^{(r)})$ is computed or estimated for each child node. The branching score of part j is defined as

$$\text{score}(j) = \max_{r \in \mathcal{K}(N)} LB(N_j^{(r)}) - \min_{r \in \mathcal{K}(N)} LB(N_j^{(r)}).$$

The selected branching part is then

$$j^* = \arg \max_{j \in \mathcal{U}(N)} \text{score}(j).$$

This rule favors parts whose machine assignment has a large impact on the node lower bound. As a result, the generated child nodes tend to have more distinct lower bounds, which improves the ability of the branch-and-bound algorithm to identify and prune inferior branches.

To reduce the computational overhead of evaluating all unassigned parts, a restricted version can also be used. In this variant, the score is computed only for a candidate subset $\mathcal{P}(N) \subseteq \mathcal{U}(N)$, such as the parts with the earliest due dates or the largest projected areas. The branching part is then selected as

$$j^* = \arg \max_{j \in \mathcal{P}(N)} \text{score}(j).$$

This restricted strong-branching strategy provides a balance between branching quality and computational efficiency. “

C. Node Lower Bound

This subsection presents the lower-bound evaluation used for partial assignment nodes. Since the upper-level search tree fixes only part-to-machine assignments, while the batching and sequencing decisions on each machine are deferred to the single-machine oracle, the lower bound is constructed at the machine-set level. The bound consists of three components: a machine-wise monotonicity bound, a memory-based bound obtained from previous oracle calls, and a fast analytical bound. An additional relaxation is further introduced to account for the unassigned parts.

1) *Machine-Wise Monotonicity Bound*: For any part subset $\mathcal{Q} \subseteq \mathcal{J}$, let $\Phi(\mathcal{Q})$ denote the optimal total tardiness of the corresponding single-machine problem in which all parts in $\mathcal{Q}$ are processed on one machine starting from time zero. The following monotonicity property is used to construct a valid lower bound for partial assignment nodes.

Proposition 1. For any two part subsets $\mathcal{Q}_1$ and $\mathcal{Q}_2$ satisfying $\mathcal{Q}_1 \subseteq \mathcal{Q}_2 \subseteq \mathcal{J}$, we have

$$\Phi(\mathcal{Q}_1) \leq \Phi(\mathcal{Q}_2).$$

Proof. Consider an optimal single-machine schedule for $\mathcal{Q}_2$. By removing all parts in $\mathcal{Q}_2 \setminus \mathcal{Q}_1$ from this schedule and keeping the relative order and batch membership of the remaining parts, we obtain a feasible schedule for $\mathcal{Q}_1$. Removing parts cannot increase the processing time of any remaining batch, because the volume component is reduced and the maximum height of a batch cannot increase. Therefore, the completion time of each part in $\mathcal{Q}_1$ does not increase. Since part tardiness is nondecreasing in its completion time, the total tardiness of the remaining parts is no larger than their contribution in the original schedule for $\mathcal{Q}_2$. Moreover, the parts in $\mathcal{Q}_2 \setminus \mathcal{Q}_1$ have nonnegative

tardiness. Hence, the optimal value for $\mathcal{Q}_1$ cannot exceed that for $\mathcal{Q}_2$, which proves the proposition. $\square$

For a partial assignment node N , the set $\mathcal{S}_m(N)$ contains the parts that have already been fixed to machine m . In any descendant of node N , machine m must process at least these parts, although additional unassigned parts may later be assigned to the same machine. By Proposition 1, the optimal total tardiness contribution of machine m in any completion of node N is at least $\Phi(\mathcal{S}_m(N))$. Therefore, the following basic node lower bound is valid:

$$LB^{\text{mono}}(N) = \sum_{m \in \mathcal{M}} \Phi(\mathcal{S}_m(N)).$$

This bound is strong when the exact values $\Phi(\mathcal{S}_m(N))$ are available. However, computing them exactly at every node would require a large number of calls to the single-machine oracle. Therefore, practical lower bounds are introduced below.

2) *Memory-Based Lower Bound*: During the search, some single-machine subproblems are solved exactly by the oracle, especially when complete assignments are evaluated or when a machine-wise subset is selected for exact evaluation. These solved subsets and their optimal values are stored in a memory pool. Let $\mathcal{C}$ denote the collection of part subsets whose exact single-machine optimal values have been computed. For each $\mathcal{Q} \in \mathcal{C}$, the value $\Phi(\mathcal{Q})$ is stored.

For a given assigned set $\mathcal{S}_m(N)$, the memory-based lower bound is defined as

$$LB_m^{\text{mem}}(\mathcal{S}_m(N)) = \max \{ \Phi(\mathcal{Q}) : \mathcal{Q} \in \mathcal{C}, \mathcal{Q} \subseteq \mathcal{S}_m(N) \}.$$

If no stored subset is contained in $\mathcal{S}_m(N)$, the value of $LB_m^{\text{mem}}(\mathcal{S}_m(N))$ is set to zero.

The validity of this bound follows directly from Proposition 1. For any stored subset $\mathcal{Q} \subseteq \mathcal{S}_m(N)$, we have

$$\Phi(\mathcal{Q}) \leq \Phi(\mathcal{S}_m(N)).$$

Taking the maximum over all such stored subsets therefore yields a valid lower bound on $\Phi(\mathcal{S}_m(N))$. The corresponding memory-based node lower bound is

$$LB^{\text{mem}}(N) = \sum_{m \in \mathcal{M}} LB_m^{\text{mem}}(\mathcal{S}_m(N)).$$

This mechanism allows the algorithm to reuse information obtained from previous oracle calls. As the search proceeds, the memory pool becomes richer, and the memory-based bound can become increasingly effective for pruning nodes whose assigned sets contain previously solved difficult subsets.

3) *Fast Analytical Lower Bound*: In addition to the memory-based bound, we compute a fast analytical lower bound for each machine-wise assigned set. For a node N and a machine $m \in \mathcal{M}$, let

$$\mathcal{Q}_m(N) = \mathcal{S}_m(N)$$

denote the set of parts that have already been assigned to machine m . The purpose of this bound is to estimate

$\Phi(\mathcal{Q}_m(N))$ without invoking the exact single-machine oracle.

If $\mathcal{Q}_m(N) = \emptyset$, we set

$$LB_m^{\text{fast}}(\mathcal{Q}_m(N)) = 0.$$

Otherwise, two analytical lower bounds are evaluated.

The first one is the parallel lower bound. It relaxes the single-machine sequencing restriction by assuming that each part in $\mathcal{Q}_m(N)$ can be processed independently as a singleton batch starting from time zero. Thus,

$$LB_m^{\text{par}}(\mathcal{Q}_m(N)) = \sum_{j \in \mathcal{Q}_m(N)} \max \{0, S + Vv_j + Uh_j - d_j\}.$$

The second one is the positional lower bound. Let $q_m = |\mathcal{Q}_m(N)|$. Sort the parts in $\mathcal{Q}_m(N)$ by projected area, height, volume, and due date, respectively:

$$\begin{aligned} a_{(1)} &\leq \dots \leq a_{(q_m)}, & h_{[1]} &\leq \dots \leq h_{[q_m]}, \\ v_{\langle 1 \rangle} &\leq \dots \leq v_{\langle q_m \rangle}, & d_{[1]}^\uparrow &\leq \dots \leq d_{[q_m]}^\uparrow. \end{aligned}$$

For each position $k = 1, \dots, q_m$, define

$$\begin{aligned} A_k &= \sum_{r=1}^k a_{(r)}, & \beta_k &= \left\lceil \frac{A_k}{LW} \right\rceil, \\ \underline{V}_k &= \sum_{r=1}^k v_{\langle r \rangle}, & \underline{H}_k &= \sum_{r=1}^{\beta_k-1} h_{[r]} + h_{[k]}. \end{aligned}$$

Then the completion time of the k -th completed part is lower bounded by

$$\underline{C}_k = \beta_k S + V \underline{V}_k + U \underline{H}_k.$$

The positional lower bound is computed as

$$LB_m^{\text{pos}}(\mathcal{Q}_m(N)) = \sum_{k=1}^{q_m} \max \{0, \underline{C}_k - d_{[k]}^\uparrow\}.$$

The derivation and validity proof of this bound are given in Appendix .

The fast analytical lower bound is defined as

$$LB_m^{\text{fast}}(\mathcal{Q}_m(N)) = \max \{LB_m^{\text{par}}(\mathcal{Q}_m(N)), LB_m^{\text{pos}}(\mathcal{Q}_m(N))\}.$$

Since both $LB_m^{\text{par}}(\mathcal{Q}_m(N))$ and $LB_m^{\text{pos}}(\mathcal{Q}_m(N))$ are valid lower bounds on $\Phi(\mathcal{Q}_m(N))$, their maximum is also valid:

$$LB_m^{\text{fast}}(\mathcal{Q}_m(N)) \leq \Phi(\mathcal{Q}_m(N)).$$

The machine-wise lower bound used in the assignment search is then

$$LB_m(\mathcal{Q}_m(N)) = \max \{LB_m^{\text{mem}}(\mathcal{Q}_m(N)), LB_m^{\text{fast}}(\mathcal{Q}_m(N))\}.$$

Therefore, the assigned-part lower bound of node N is

$$LB^{\text{ass}}(N) = \sum_{m \in \mathcal{M}} LB_m(\mathcal{Q}_m(N)).$$

4) *Unassigned-Part Relaxation:* The bound $LB^{\text{ass}}(N)$ accounts only for the parts that have already been assigned to machines. At early nodes, this bound can be weak because many parts remain unassigned. To strengthen the node evaluation, we add a relaxation for the unassigned parts.

For each unassigned part $j \in \mathcal{U}(N)$, its completion time in any feasible schedule cannot be earlier than the processing time of a singleton batch containing only part j and starting at time zero. This singleton processing time is

$$p_j^{\text{sing}} = S + Vv_j + Uh_j.$$

Therefore, the tardiness contribution of part j is at least

$$\max \{0, p_j^{\text{sing}} - d_j\}$$

under this relaxation. Summing over all unassigned parts yields

$$LB^{\text{U}}(N) = \sum_{j \in \mathcal{U}(N)} \max \{0, S + Vv_j + Uh_j - d_j\}.$$

This is a valid lower bound on the total tardiness contribution of the unassigned parts. The relaxation is optimistic because it allows every unassigned part to be processed independently as a singleton batch starting from time zero, without considering machine capacity, machine occupation, or competition among unassigned parts.

Combining the assigned-part bound and the unassigned-part relaxation, the final node lower bound is defined as

$$LB(N) = LB^{\text{ass}}(N) + LB^{\text{U}}(N).$$

Equivalently,

$$LB(N) = \sum_{m \in \mathcal{M}} LB_m(\mathcal{Q}_m(N)) + \quad (15)$$

$$\sum_{j \in \mathcal{U}(N)} \max \{0, S + Vv_j + Uh_j - d_j\}. \quad (16)$$

The validity of $LB(N)$ follows from the fact that the first term is a lower bound on the tardiness contribution of all parts already assigned to machines, while the second term is a lower bound on the tardiness contribution of the remaining unassigned parts. Since these two sets of parts are disjoint, their lower bounds can be added. A node N can therefore be safely pruned whenever

$$LB(N) \geq UB,$$

where UB is the incumbent upper bound.

D. Upper Bound and Complete Assignment Evaluation

An incumbent upper bound is required to prune nodes during the branch-and-bound search. In the proposed framework, an upper bound is obtained from any complete feasible assignment of parts to machines, because the batching and sequencing decisions on each machine can then be evaluated exactly by the single-machine oracle.

1) *Initial Upper Bound*: At the beginning of the search, a constructive heuristic is used to generate an initial complete assignment. The heuristic first sorts all parts in nondecreasing order of due dates. The parts are then assigned one by one to the parallel machines. When considering a part j , the heuristic evaluates the effect of assigning j to each candidate machine and selects the machine that yields the smallest estimated increase in the total tardiness.

Let $\hat{\Phi}(\mathcal{Q})$ denote a fast heuristic estimate of the single-machine total tardiness for processing the part set $\mathcal{Q}$ on one machine. For the current partial assignment $\{\mathcal{S}_m\}_{m \in \mathcal{M}}$, the selected machine for part j is given by

$$r^* = \arg \min_{r \in \mathcal{M}} \left\{ \hat{\Phi}(\mathcal{S}_r \cup \{j\}) - \hat{\Phi}(\mathcal{S}_r) \right\}.$$

Part j is then assigned to machine r^* , and the assigned set of that machine is updated as

$$\mathcal{S}_{r^*} \leftarrow \mathcal{S}_{r^*} \cup \{j\}.$$

After all parts have been assigned, the resulting complete assignment is evaluated exactly. For each machine $m \in \mathcal{M}$, the exact single-machine oracle is called to compute $\Phi(\mathcal{S}_m)$. The initial upper bound is then set as

$$UB = \sum_{m \in \mathcal{M}} \Phi(\mathcal{S}_m).$$

The corresponding assignment and the associated single-machine schedules are stored as the initial incumbent solution.

In practice, several constructive rules can be used to generate multiple initial assignments, such as earliest due date, largest projected area, largest volume, and largest height. Each complete assignment is evaluated exactly by the oracle, and the best one is selected as the initial incumbent. This strategy increases the chance of obtaining a tight initial upper bound while requiring only a limited number of oracle calls.

2) *Complete Assignment Evaluation*: A node N is a complete assignment if

$$\mathcal{U}(N) = \emptyset.$$

At such a node, every part has been assigned to exactly one machine. Therefore, the remaining optimization decisions are independent across machines. For each machine m , the assigned part set $\mathcal{S}_m(N)$ defines a single-machine AM batch scheduling subproblem.

The exact value of the complete assignment is computed as

$$Z(N) = \sum_{m \in \mathcal{M}} \Phi(\mathcal{S}_m(N)),$$

where $\Phi(\mathcal{S}_m(N))$ is obtained by the exact single-machine oracle. If $\mathcal{S}_m(N) = \emptyset$, then

$$\Phi(\mathcal{S}_m(N)) = 0.$$

To avoid repeated computations, all oracle evaluations are stored in the memory pool $\mathcal{C}$. Before calling the oracle for a set $\mathcal{S}_m(N)$, the algorithm first checks whether this set has already been evaluated. If $\mathcal{S}_m(N) \in \mathcal{C}$, the stored value

$\Phi(\mathcal{S}_m(N))$ is directly retrieved. Otherwise, the oracle is called, and the computed value is added to the memory pool.

Once $Z(N)$ is obtained, it is compared with the current incumbent upper bound. If

$$Z(N) < UB,$$

then the incumbent solution is updated and the upper bound is reset as

$$UB \leftarrow Z(N).$$

Otherwise, the complete assignment is discarded.

Because each leaf node is evaluated by an exact single-machine oracle, the value $Z(N)$ is the true objective value of the corresponding complete parallel-machine assignment. Consequently, the incumbent solution maintained by the algorithm is always feasible for the original problem, and its value is a valid upper bound throughout the search.

E. Node Selection and Pruning

The active nodes generated during the assignment-based search are stored in an active node list, denoted by $\mathcal{L}$. The order in which nodes are selected from $\mathcal{L}$ affects both the speed of finding high-quality incumbent solutions and the efficiency of proving optimality. Therefore, the proposed algorithm adopts a hybrid node selection strategy that combines depth-first and best-first search.

1) *Node Selection Strategy*: A best-first strategy selects the active node with the smallest lower bound:

$$N^* = \arg \min_{N \in \mathcal{L}} LB(N).$$

This strategy is useful for improving the global lower bound and proving optimality. However, in the early stage of the search, best-first selection may delay the discovery of complete assignments and may therefore provide weak incumbent upper bounds.

A depth-first strategy, by contrast, selects one of the deepest active nodes in the search tree. Let $d(N)$ denote the depth of node N , namely the number of parts that have already been assigned. A depth-first node can be selected as

$$N^* = \arg \max_{N \in \mathcal{L}} d(N).$$

This strategy quickly generates complete assignments and can therefore help obtain high-quality upper bounds at an early stage. However, it may also spend excessive time exploring unpromising regions of the search tree.

To balance these two effects, we use a hybrid node selection strategy. Although an initial incumbent is available before the search starts, the algorithm first uses depth-first search for a short phase to improve the incumbent quickly. It then switches to best-first search to accelerate the proof of optimality. In addition, two thresholds, $N_{\max}$ and $N_{\min}$, are used to control the size of the active node list. If the number of active nodes exceeds $N_{\max}$, the algorithm switches to depth-first mode to reduce memory consumption. Once the number of active nodes decreases below $N_{\min}$, the algorithm switches back to best-first mode.

More specifically, the node selection rule is defined as follows:

$$N^* = \begin{cases} \arg \max_{N \in \mathcal{L}} d(N), & \text{if depth-first mode is active,} \\ \arg \min_{N \in \mathcal{L}} LB(N), & \text{if best-first mode is active.} \end{cases}$$

When multiple nodes have the same selection priority, ties are broken by choosing the node with the larger depth. This tie-breaking rule encourages the search to reach complete assignments earlier, while still following the main priority rule.

2) *Pruning Rules*: Several pruning rules are applied to reduce the search space.

Bound-based pruning.: Let UB denote the current incumbent upper bound. For any active node N , if

$$LB(N) \geq UB,$$

then no descendant of N can improve the incumbent solution. Therefore, node N is safely pruned.

Terminal-node evaluation.: If node N is terminal, namely

$$\mathcal{U}(N) = \emptyset,$$

then all parts have been assigned to machines. The node is evaluated exactly by calling the single-machine oracle for each machine:

$$Z(N) = \sum_{m \in \mathcal{M}} \Phi(\mathcal{S}_m(N)).$$

If $Z(N) < UB$, the incumbent solution is updated and UB is set to $Z(N)$. After this evaluation, the terminal node is discarded because it has no children.

Cache-based pruning and reuse.: Before evaluating a machine-wise subset $\mathcal{S}_m(N)$ by the exact single-machine oracle, the algorithm checks whether this subset has already been solved and stored in the memory pool $\mathcal{C}$. If $\mathcal{S}_m(N) \in \mathcal{C}$, the stored value $\Phi(\mathcal{S}_m(N))$ is retrieved directly. Otherwise, the oracle is called and the resulting value is added to $\mathcal{C}$. This mechanism does not alter the search space, but it avoids repeated exact evaluations and strengthens the memory-based lower bound for subsequent nodes.

Symmetry-based pruning.: The symmetry-reduced branching scheme described in Section II-B prevents the generation of nodes that differ only by a relabeling of identical machines. In particular, a new machine can be opened only after all earlier machine indices have already been opened. As a result, symmetric assignment patterns are removed before entering the active node list.

3) *Optimality Guarantee*: The proposed pruning rules do not remove any node that may contain a solution better than the incumbent. Bound-based pruning is valid because $LB(N)$ is a lower bound on the objective value of all complete assignments descending from node N . Terminal nodes are evaluated exactly using the single-machine oracle. The symmetry-reduced branching rule preserves at least one representative for every class of equivalent assignments under machine relabeling. Therefore, when the active node list becomes empty, all nonpruned complete assignments have

been evaluated and all pruned nodes have been proven unable to improve the incumbent. The final incumbent solution is therefore globally optimal.

F. Exact Single-Machine Scheduling Oracle

In the proposed assignment-based branch-and-bound framework, each terminal node fixes a complete assignment of parts to parallel machines. Once the assignment is fixed, the problem decomposes into independent single-machine AM batch scheduling subproblems. To evaluate each subproblem exactly, we use a dedicated single-machine scheduling oracle.

For a given part subset $\mathcal{Q} \subseteq \mathcal{J}$, the oracle solves the following problem: determine the batching and sequencing decisions on a single AM machine so as to minimize the total tardiness of the parts in $\mathcal{Q}$. The optimal objective value returned by the oracle is denoted by

$$\Phi(\mathcal{Q}).$$

If $\mathcal{Q} = \emptyset$, we set $\Phi(\mathcal{Q}) = 0$.

The oracle is implemented as an exact branch-and-bound algorithm. Each node represents a partial batch sequence fixed from time zero and a set of remaining unscheduled parts. Child nodes are generated by selecting a feasible subset of the remaining parts as the next batch. The search is accelerated by a constructive upper bound, valid lower bounds, adaptive node selection, and dominance-based pruning. Since the branching scheme enumerates all feasible batch sequences and pruning is performed only by valid lower bounds and dominance rules, the oracle returns the globally optimal value $\Phi(\mathcal{Q})$ whenever it terminates with an empty active node list.

All oracle calls are cached. Therefore, if the same subset $\mathcal{Q}$ is encountered again during the parallel-machine search, the stored value $\Phi(\mathcal{Q})$ is retrieved directly without resolving the subproblem. The detailed oracle implementation is provided in Appendix .

III. TEST INSTANCES

The computational experiments are conducted on test instances adopted from Mao et al. [3]. For each part j , the due date d_j is generated following the procedure of Yu et al. [10]. Specifically, due dates are sampled from a uniform distribution:

$$d_j \sim \text{Unif}[\mu(1 - TF - RDD/2), \mu(1 - TF + RDD/2)] \quad (17)$$

where TF and RDD denote the tardiness factor and the relative due-date range, respectively. A larger value of TF corresponds to tighter due dates and therefore a more time-critical demand setting. The parameter μ represents an estimate of the makespan of the instance. For the single-machine setting considered here, μ is estimated by a simple first-fit heuristic in which parts are sorted in non-decreasing order of projected area. In these instances, both part dimensions range from 10 to 19, and the machine platform size is 20×20 .

IV. CONCLUSIONS

REFERENCES

- [1] M. Pinto, C. Silva, M. Thürrer, and S. Moniz, "Nesting and scheduling optimization of additive manufacturing systems: Mapping the territory," *Computers & Operations Research*, vol. 165, p. 106592, May 2024.
- [2] H. Wu, C. Yu, and S. Yu, "Nesting and scheduling for parallel additive manufacturing machines with uncertain processing times: A simulation-optimisation approach," *International Journal of Production Research*, pp. 1–30, Jun. 2025.
- [3] Z. Mao, E. Fu, D. Huang, K. Fang, and L. Chen, "Combinatorial Benders decomposition for single machine scheduling in additive manufacturing with two-dimensional packing constraints," *European Journal of Operational Research*, vol. 317, no. 3, pp. 890–905, Sep. 2024.
- [4] B. Zipfel, F. Tamke, and L. Kuttner, "A new branch-and-cut approach for integrated planning in additive manufacturing," *European Journal of Operational Research*, p. S0377221724008439, Nov. 2024.
- [5] P. J. Nascimento, C. Silva, C. H. Antunes, and S. Moniz, "Optimal decomposition approach for solving large nesting and scheduling problems of additive manufacturing systems," *European Journal of Operational Research*, vol. 317, no. 1, pp. 92–110, Aug. 2024.
- [6] P. J. Nascimento, C. Silva, C. H. Antunes, C. T. Maravelias, and S. Moniz, "Improving the efficiency of Logic-based Benders Decomposition for p-batch scheduling problems with two-dimensional packing," *European Journal of Operational Research*, p. S0377221726000767, Jan. 2026.
- [7] Y. Che, K. Hu, Z. Zhang, and A. Lim, "Machine scheduling with orientation selection and two-dimensional packing for additive manufacturing," *Computers & Operations Research*, vol. 130, p. 105245, Jun. 2021.
- [8] S. K. Tangudu and M. E. Kurz, "A branch and bound algorithm to minimise total weighted tardiness on a single batch processing machine with ready times and incompatible job families," *Production Planning & Control*, vol. 17, no. 7, pp. 728–741, Oct. 2006.
- [9] E. L. Lawler, "A "Pseudopolynomial" Algorithm for Sequencing Jobs to Minimize Total Tardiness," in *Annals of Discrete Mathematics*. Elsevier, 1977, vol. 1, pp. 331–342.
- [10] C. Yu, A. Matta, Q. Semeraro, and J. Lin, "Mathematical Models for Minimizing Total Tardiness on Parallel Additive Manufacturing Machines," *IFAC-PapersOnLine*, vol. 55, no. 10, pp. 1521–1526, 2022.

APPENDIX

This appendix derives the positional lower bound used in Section II-C.3. Consider an arbitrary part set $\mathcal{Q} \subseteq \mathcal{J}$ to be processed on a single machine starting from time zero. Let $q = |\mathcal{Q}|$. For each part $j \in \mathcal{Q}$, let

$$a_j = l_j w_j$$

denote its projected area. Sort the parts in $\mathcal{Q}$ by projected area, height, and volume:

$$a_{(1)} \leq a_{(2)} \leq \dots \leq a_{(q)},$$

$$h_{[1]} \leq h_{[2]} \leq \dots \leq h_{[q]},$$

$$v_{\langle 1 \rangle} \leq v_{\langle 2 \rangle} \leq \dots \leq v_{\langle q \rangle}.$$

For each position $k = 1, \dots, q$, define

$$A_k = \sum_{r=1}^k a_{(r)}, \quad \beta_k = \left\lceil \frac{A_k}{LW} \right\rceil,$$

and

$$\underline{V}_k = \sum_{r=1}^k v_{\langle r \rangle}.$$

The quantity A_k is the minimum possible total projected area of any k parts in $\mathcal{Q}$. Therefore, before the k -th part can

be completed, at least β_k batches must have been processed. This gives a setup-time lower bound of $\beta_k S$. Similarly, the total processed volume before the k -th completion must be at least the sum of the k smallest volumes, namely $\underline{V}_k$. Hence, the volume-dependent processing-time contribution is at least $V \underline{V}_k$.

It remains to lower bound the cumulative height-dependent contribution. Consider the first k completed parts in any feasible schedule. These parts must be distributed over at least β_k nonempty batches before the k -th completion. Let H_b denote the maximum height of batch b . The cumulative height contribution of these batches is therefore at least the optimal value of the following auxiliary problem:

$$H_k^* = \min \sum_{b=1}^r H_b$$

subject to

$$B_b \subseteq \{1, \dots, k\}, \quad b = 1, \dots, r,$$

$$\bigcup_{b=1}^r B_b = \{1, \dots, k\},$$

$$B_b \cap B_{b'} = \emptyset, \quad 1 \leq b < b' \leq r,$$

$$B_b \neq \emptyset, \quad b = 1, \dots, r,$$

$$\sum_{j \in B_b} a_j \leq LW, \quad b = 1, \dots, r,$$

$$H_b = \max_{j \in B_b} h_j, \quad b = 1, \dots, r,$$

$$r \geq \beta_k.$$

This auxiliary problem is still combinatorial. We therefore relax it by fixing the number of groups to β_k and removing the capacity constraints. The relaxed problem is

$$\tilde{H}_k = \min \sum_{b=1}^{\beta_k} H_b$$

subject to

$$G_b \subseteq \{1, \dots, k\}, \quad b = 1, \dots, \beta_k,$$

$$\bigcup_{b=1}^{\beta_k} G_b = \{1, \dots, k\},$$

$$G_b \cap G_{b'} = \emptyset, \quad 1 \leq b < b' \leq \beta_k,$$

$$G_b \neq \emptyset, \quad b = 1, \dots, \beta_k,$$

$$H_b = \max_{j \in G_b} h_{[j]}, \quad b = 1, \dots, \beta_k.$$

By construction, this is a relaxation of the auxiliary height problem. Therefore,

$$\tilde{H}_k \leq H_k^*.$$

Lemma 1.: For each $k = 1, \dots, q$, the optimal value of the relaxed height problem is

$$\tilde{H}_k = \sum_{r=1}^{\beta_k-1} h_{[r]} + h_{[k]}.$$

Proof.: Since the β_k groups must be nonempty, each group contributes at least the height of one assigned part through its group maximum. To minimize the sum of group maxima, $\beta_k - 1$ groups should be made singleton groups containing the $\beta_k - 1$ smallest heights $h_{[1]}, \dots, h_{[\beta_k-1]}$. All remaining parts are assigned to the last group, whose maximum height is $h_{[k]}$. Therefore,

$$\tilde{H}_k = \sum_{r=1}^{\beta_k-1} h_{[r]} + h_{[k]}.$$

This proves the lemma. $\square$

Define

$$\underline{H}_k = \tilde{H}_k = \sum_{r=1}^{\beta_k-1} h_{[r]} + h_{[k]}.$$

Combining the setup, volume, and height contributions, define

$$\underline{C}_k = \beta_k S + V \underline{V}_k + U \underline{H}_k, \quad k = 1, \dots, q.$$

Proposition 2.: Let $C_{(k)}^{\text{real}}$ denote the completion time of the k -th completed part among the parts in $\mathcal{Q}$ in any feasible single-machine schedule. Then,

$$C_{(k)}^{\text{real}} \geq \underline{C}_k, \quad k = 1, \dots, q.$$

Proof.: Fix any $k \in \{1, \dots, q\}$ and consider the first k completed parts in an arbitrary feasible schedule.

First, these k parts have total projected area at least

$$A_k = \sum_{r=1}^k a_{(r)}.$$

Since each batch has capacity LW , at least

$$\beta_k = \left\lceil \frac{A_k}{LW} \right\rceil$$

batches must have been processed before the k -th completion. Hence, the cumulative setup contribution is at least $\beta_k S$.

Second, the total processed volume before the k -th completion must be at least the sum of the k smallest volumes in $\mathcal{Q}$:

$$\underline{V}_k = \sum_{r=1}^k v_{(r)}.$$

Thus, the cumulative volume-dependent contribution is at least $V \underline{V}_k$.

Third, the cumulative height-dependent contribution before the k -th completion is at least $U \underline{H}_k^*$. Since

$$H_k^* \geq \tilde{H}_k = \underline{H}_k,$$

this contribution is further bounded below by $U \underline{H}_k$.

Combining the three components gives

$$C_{(k)}^{\text{real}} \geq \beta_k S + V \underline{V}_k + U \underline{H}_k = \underline{C}_k.$$

This proves the proposition. $\square$

To transform the completion-time lower bounds into a total tardiness lower bound, sort the due dates of the parts in $\mathcal{Q}$ in nondecreasing order:

$$d_{[1]}^\uparrow \leq d_{[2]}^\uparrow \leq \dots \leq d_{[q]}^\uparrow.$$

Since the tardiness function is nondecreasing in completion time, the minimum tardiness implied by the positional completion-time lower bounds is obtained by matching smaller completion-time lower bounds with smaller due dates. Hence,

$$LB^{\text{pos}}(\mathcal{Q}) = \sum_{k=1}^q \max \left\{ 0, \underline{C}_k - d_{[k]}^\uparrow \right\}$$

is a valid lower bound on the optimal single-machine total tardiness $\Phi(\mathcal{Q})$.

Applying this expression to $\mathcal{Q} = \mathcal{Q}_m(N) = \mathcal{S}_m(N)$ yields the machine-wise positional lower bound $LB_m^{\text{pos}}(\mathcal{Q}_m(N))$ used in the main algorithm.

APPENDIX

This appendix describes the exact single-machine scheduling oracle used in the parallel-machine branch-and-bound algorithm. For a given part subset $\mathcal{Q} \subseteq \mathcal{J}$, the oracle solves the single-machine one-dimensional AM batch scheduling problem with total tardiness minimization. The optimal objective value returned by the oracle is denoted by $\Phi(\mathcal{Q})$.

A. Single-Machine Subproblem

Given a part subset $\mathcal{Q}$, the single-machine subproblem is to partition the parts in $\mathcal{Q}$ into feasible batches and sequence these batches on one AM machine. A batch $B \subseteq \mathcal{Q}$ is feasible if

$$\sum_{j \in B} l_j w_j \leq LW. \quad (18)$$

The processing time of a feasible batch B is

$$p(B) = S + V \sum_{j \in B} v_j + U \max_{j \in B} h_j. \quad (19)$$

If batch B is completed at time $C(B)$, then every part $j \in B$ has completion time $C(B)$ and tardiness

$$T_j = \max\{0, C(B) - d_j\}. \quad (20)$$

The oracle returns

$$\Phi(\mathcal{Q}) = \min \sum_{j \in \mathcal{Q}} T_j. \quad (21)$$

If $\mathcal{Q} = \emptyset$, we define $\Phi(\mathcal{Q}) = 0$.

B. Node Representation

The oracle is implemented by a branch-and-bound algorithm that builds the batch sequence from time zero. Each node N is represented by

$$N = (\mathcal{B}(N), \mathcal{U}(N), C_N, TT_N), \quad (22)$$

where $\mathcal{B}(N)$ is the sequence of batches already fixed, $\mathcal{U}(N)$ is the set of unscheduled parts, C_N is the completion time

of the last fixed batch, and TT_N is the total tardiness contributed by the already scheduled parts.

The root node is

$$N^0 = (\emptyset, \mathcal{Q}, 0, 0). \quad (23)$$

A node is terminal if

$$\mathcal{U}(N) = \emptyset. \quad (24)$$

At a terminal node, TT_N is the total tardiness of a complete feasible single-machine schedule.

C. Branching Rule

At a nonterminal node N , child nodes are generated by selecting a nonempty feasible subset $B \subseteq \mathcal{U}(N)$ as the next batch. The subset must satisfy

$$\sum_{j \in B} l_j w_j \leq LW. \quad (25)$$

For each feasible subset B , a child node N' is generated as follows:

$$\mathcal{B}(N') = \mathcal{B}(N) \oplus B, \quad (26)$$

where $\oplus$ denotes appending B to the current batch sequence,

$$\mathcal{U}(N') = \mathcal{U}(N) \setminus B, \quad (27)$$

$$C_{N'} = C_N + p(B), \quad (28)$$

and

$$TT_{N'} = TT_N + \sum_{j \in B} \max\{0, C_{N'} - d_j\}. \quad (29)$$

This branching scheme is complete because any feasible single-machine batch schedule can be uniquely represented as a sequence of feasible batches selected from the remaining parts at each stage.

D. Initial Upper Bound

Before the branch-and-bound search starts, an initial feasible solution is constructed by a fast heuristic. The parts in $\mathcal{Q}$ are first sorted in nondecreasing order of due dates. The heuristic then scans the sorted list and assigns each part to the current batch if the platform capacity constraint remains satisfied. Otherwise, the current batch is closed and a new batch is opened. The resulting batch sequence is evaluated exactly by (19), and its total tardiness is used as the initial upper bound UB .

In the implementation, several simple constructive rules can be used, such as earliest due date, largest projected area, largest volume, and largest height. The best schedule obtained from these rules is selected as the initial incumbent.

E. Parallel Lower Bound

At a node N , the parallel lower bound relaxes the single-machine restriction by assuming that each unscheduled part can be processed independently as a singleton batch immediately after the current completion time C_N . For each part $j \in \mathcal{U}(N)$, the earliest completion time under this relaxation is

$$C_N + S + Vv_j + Uh_j. \quad (30)$$

Therefore, the parallel lower bound is defined as

$$LB_{\text{par}}(N) = TT_N + \sum_{j \in \mathcal{U}(N)} \max\{0, C_N + S + Vv_j + Uh_j - d_j\}. \quad (31)$$

This bound is valid because no feasible continuation can complete part j before processing at least its own volume and height contribution together with one setup operation after time C_N .

F. Serial Lower Bound

Although $LB_{\text{par}}(N)$ is inexpensive, it ignores the sequential nature of the remaining decisions. A stronger serial lower bound is therefore constructed.

Consider an arbitrary feasible continuation from node N . Suppose that the unscheduled parts are partitioned into batches

$$B_1, B_2, \dots, B_m$$

and these batches are processed sequentially after time C_N . For a part $j \in B_k$, its actual completion time is

$$C_j^{\text{real}} = C_N + \sum_{r=1}^k \left(S + V \sum_{i \in B_r} v_i + U \max_{i \in B_r} h_i \right). \quad (32)$$

Let π be the part sequence induced by listing the parts batch by batch according to the order $B_1, B_2, \dots, B_m$. Let

$$h_{\min} = \min_{j \in \mathcal{U}(N)} h_j. \quad (33)$$

A relaxed completion time for the part at position p in π is defined as

$$\tilde{C}_{\pi(p)} = C_N + S + Uh_{\min} + V \sum_{r=1}^p v_{\pi(r)}. \quad (34)$$

This relaxation charges the setup term and the height-dependent term only once for all remaining parts, while preserving the cumulative volume-dependent contribution along the induced sequence.

For a sequence π , let

$$T_{\text{relax}}(\pi) = \sum_{p=1}^{|\mathcal{U}(N)|} \max\{0, \tilde{C}_{\pi(p)} - d_{\pi(p)}\}. \quad (35)$$

Let π^* be an optimal sequence for the relaxed problem. The serial lower bound is then

$$LB_{\text{ser}}(N) = TT_N + T_{\text{relax}}(\pi^*). \quad (36)$$

The relaxed problem is equivalent to a single-machine total tardiness problem $1||\sum T_j$ over the unscheduled parts,

where part j has processing time Vv_j and all parts share the same offset $C_N + S + Uh_{\min}$. In the implementation, this relaxed problem is solved exactly by a dynamic programming procedure.

Proposition 1. *For any node N , $LB_{\text{ser}}(N)$ is a valid lower bound on the total tardiness of any complete schedule extending N .*

Proof. Consider any feasible continuation from node N and its induced sequence π . For the part $\pi(p)$, suppose that it belongs to batch B_k . Its actual completion time in (32) includes all setup, volume, and height contributions of batches $B_1, \dots, B_k$. Therefore, it contains at least one setup term, at least one height-dependent term no smaller than $Uh_{\min}$, and a cumulative volume contribution no smaller than that of the first p parts in the induced sequence. Hence,

$$C_{\pi(p)}^{\text{real}} \geq \tilde{C}_{\pi(p)}, \quad p = 1, \dots, |\mathcal{U}(N)|. \quad (37)$$

Since the tardiness function is nondecreasing in completion time, the total tardiness of the feasible continuation is no smaller than $T_{\text{relax}}(\pi)$. Moreover, π^* is optimal for the relaxed problem, so

$$T_{\text{relax}}(\pi) \geq T_{\text{relax}}(\pi^*). \quad (38)$$

Adding the already fixed tardiness TT_N gives

$$TT_N + T_{\text{batch}}(\pi) \geq TT_N + T_{\text{relax}}(\pi^*) = LB_{\text{ser}}(N). \quad (39)$$

Thus, $LB_{\text{ser}}(N)$ is a valid lower bound. $\square$

G. Hybrid Lower Bound

The parallel lower bound is fast but relatively weak, whereas the serial lower bound is stronger but more expensive. To balance pruning strength and computational effort, a hybrid lower bound is used.

At each node, $LB_{\text{par}}(N)$ is computed first. The serial lower bound is evaluated only when the node is close to being pruned and still contains a sufficiently large number of unscheduled parts. Specifically, the hybrid lower bound is defined as

$$LB_{\text{hyb}}(N) = \begin{cases} \max\{LB_{\text{par}}(N), LB_{\text{ser}}(N)\}, & \text{if } \frac{UB - LB_{\text{par}}(N)}{UB} \leq \theta_1 \text{ and } \frac{|\mathcal{U}(N)|}{|\mathcal{Q}|} > \theta_2, \\ LB_{\text{par}}(N), & \text{otherwise,} \end{cases} \quad (40)$$

where θ_1 and θ_2 are control parameters. In the computational experiments, we set $\theta_1 = 0.10$ and $\theta_2 = 0.60$.

Since both $LB_{\text{par}}(N)$ and $LB_{\text{ser}}(N)$ are valid lower bounds, their maximum is also valid. When the serial lower bound is not activated, $LB_{\text{hyb}}(N)$ reduces to the valid parallel lower bound. Therefore, $LB_{\text{hyb}}(N)$ is valid for all nodes.

H. Dominance Rule

A dominance rule is used to remove redundant partial schedules. Let u and v be two nodes. Let $\mathcal{S}_u$ and $\mathcal{S}_v$ denote the sets of already scheduled parts at these two nodes. Node u dominates node v , denoted by $u \prec v$, if

$$\mathcal{S}_u = \mathcal{S}_v, \quad (41)$$

$$C_u \leq C_v, \quad (42)$$

and

$$TT_u \leq TT_v, \quad (43)$$

with at least one inequality being strict.

Proposition 2. *If node u dominates node v , then node v can be safely pruned.*

Proof. Since $\mathcal{S}_u = \mathcal{S}_v$, the two nodes have the same unscheduled part set and admit the same feasible continuations. Consider any feasible continuation of node v . Applying the same continuation to node u is also feasible. The appended batches have identical processing times in both continuations. Since $C_u \leq C_v$, every future batch completion time from node u is no later than the corresponding completion time from node v . Thus, the future tardiness generated from node u is no larger than that generated from node v . Together with $TT_u \leq TT_v$, the total tardiness of the complete schedule obtained from node u is no larger than that obtained from node v . Therefore, node v cannot lead to a better solution than node u and can be safely pruned. $\square$

In the implementation, nondominated labels are maintained for each scheduled part set. A newly generated node is discarded if it is dominated by an existing label with the same scheduled set. Otherwise, it is inserted into the active node list, and any existing labels dominated by the new node are removed.

I. Node Selection

The active nodes are stored in a list $\mathcal{L}$. The oracle supports three node selection strategies.

In depth-first search, the selected node is one of the deepest active nodes:

$$N^* = \arg \max_{N \in \mathcal{L}} |\mathcal{B}(N)|. \quad (44)$$

In best-first search, the selected node has the smallest lower bound:

$$N^* = \arg \min_{N \in \mathcal{L}} LB_{\text{hyb}}(N). \quad (45)$$

The adaptive strategy switches between depth-first and best-first search. It uses best-first search by default. If the number of active nodes exceeds a threshold $N_{\max}$, it switches to depth-first search to reduce memory usage. Once the number of active nodes decreases below $N_{\min}$, it switches back to best-first search.

J. Overall Procedure

The exact oracle proceeds as follows.

- 1) Construct an initial feasible batch sequence and set its total tardiness as the initial upper bound UB .
- 2) Initialize the root node

$$N^0 = (\emptyset, \mathcal{Q}, 0, 0)$$

and insert it into the active node list $\mathcal{L}$.

- 3) While $\mathcal{L}$ is not empty, select a node N according to the node selection strategy.
- 4) Compute $LB_{\text{hyb}}(N)$. If

$$LB_{\text{hyb}}(N) \geq UB,$$

prune node N .

- 5) If $\mathcal{U}(N) = \emptyset$, node N is terminal. If

$$TT_N < UB,$$

update the incumbent schedule and set $UB \leftarrow TT_N$. Then discard node N .

- 6) If $\mathcal{U}(N) \neq \emptyset$, enumerate all nonempty feasible subsets $B \subseteq \mathcal{U}(N)$ satisfying

$$\sum_{j \in B} l_j w_j \leq LW.$$

For each feasible subset, generate the corresponding child node.

- 7) Apply the dominance rule to each generated child node. If the child node is dominated, discard it. Otherwise, insert it into $\mathcal{L}$ and update the dominance labels.
- 8) When $\mathcal{L}$ becomes empty, return the incumbent schedule and its objective value UB .

K. Optimality Guarantee

The oracle enumerates all feasible batch sequences through the branching scheme. A node is pruned only when its lower bound is no smaller than the current incumbent upper bound or when it is dominated by another node that has scheduled the same part set with no larger completion time and no larger accumulated tardiness. Both pruning mechanisms are valid. Therefore, when the active node list becomes empty, every feasible schedule has either been evaluated explicitly or proven unable to improve the incumbent. The returned incumbent schedule is thus globally optimal, and its objective value is $\Phi(\mathcal{Q})$.

If the oracle is terminated by a time limit before the active node list becomes empty, the returned incumbent remains feasible, but global optimality is not guaranteed.